

Job Recommendation System

Team Digital Nomads -

Chaitanya Gawande

Rutik Darda

Sagar Patel

Utkarsh Shah

M.S. Software Engineering

San Jose State University

Prof. Vijay Eranti

12/15/2023

Table of Contents

1. Abstract
2. Introduction
3. Related Work
4. Data
5. Methods
6. Experiments and Results
7. Conclusion
8. Supplementary Material

Abstract

In response to the dynamic challenges job seekers face, our data mining project focuses on optimizing the job-seeking process by developing a specialized Job Recommendation System (JRS). Traditional job search methods, often reliant on simplistic keyword matching, may overlook the nuanced skills and qualifications of candidates. In light of this, our innovative approach, firmly rooted in data mining principles, aims to revolutionize the job recommendation paradigm.

The methodology involves a meticulous extraction and analysis of skills from uploaded resumes, utilizing advanced algorithms to uncover latent connections between a candidate's skill set and specific job requirements. By incorporating natural language processing and machine learning, our JRS not only identifies explicit skills but also discerns implicit relationships and contextual nuances within textual data. The primary goal is to enhance the precision and efficiency of job recommendations.

Results from our implementation showcase the system's effectiveness in providing more accurate and relevant job suggestions compared to conventional keyword-based methods. Although we focus on recommending jobs rather than directly impacting the job market, the optimized job-seeking process benefits both candidates and recruiters. Our JRS contributes to the ongoing evolution of personalized job recommendation systems, successfully bridging the gap between job seekers and potential employers. This project underscores the capacity of data mining to elevate the job-matching experience, making it a valuable tool for individuals navigating the competitive job landscape.

This project leverages the rich LinkedIn Job Postings Kaggle Dataset to build a highly accurate and personalized job recommendation system. Initial results demonstrate high precision and recall in matching candidate profiles with relevant job postings, significantly improving job search efficiency.

Introduction

In the context of modern employment challenges, our data mining project addresses the inefficiencies in connecting job seekers with suitable opportunities using conventional keyword-based job search methods. Recognizing the limitations of these approaches, we introduce a specialized Job Recommendation System (JRS) that goes beyond simplistic matching techniques. By applying advanced algorithms, including natural language processing and machine learning, our system meticulously analyzes the inherent structure of resumes and job descriptions. The significance of this initiative lies in its potential to optimize the job-seeking process, offering more precise and relevant job suggestions for individual users. Preliminary results indicate that our JRS demonstrates superior performance compared to traditional methods, showcasing its promise in reshaping the landscape of contemporary job searching and recommendation systems. This report provides a comprehensive overview of our methodology, approach, and initial success in addressing the critical gap in the employment exploration process.

Related Work

1. GIRL - Generative Job Recommendations: GIRL focuses on enhancing the accuracy and relevance of job recommendations using advanced algorithms. Our project utilizes natural language processing and machine learning, but GIRL introduces a generative aspect, creating job descriptions based on CVs, which differs from our approach of analyzing resumes and job postings.

Paper Link - <https://arxiv.org/abs/2307.02157>

2. Job Recommendation with Progression of Applications: This paper's methodology, considering the progression of job applications, offers a dynamic aspect of job recommendations. Our project shares the use of advanced algorithms for analysis but doesn't explicitly mention considering the application history, focusing instead on matching skills and qualifications.

Paper Link- <https://arxiv.org/abs/1905.13136>

3. Job Recommender Systems Review: This comprehensive review covers various JRS types, including those similar to our project. It highlights the importance of considering the temporal and reciprocal nature of job recommendations, which is indirectly relevant to our approach of matching candidates to jobs based on skills.

Paper Link- <https://arxiv.org/abs/2111.13576>

Overall, our project shares similarities in using advanced data mining techniques and machine learning for enhanced job recommendations. However, it differs in specific methodologies, such as the generative aspect in GIRL and the consideration of application progression in the second paper.

Data

LinkedIn Job Postings Dataset:

<https://www.kaggle.com/datasets/rajatraj0502/linkedin-job-2023>

The datasets from the Kaggle Dataset for this project are:

1. Companies Dataset: This dataset contains details about various companies, including company IDs, names, descriptions, sizes, and locations (state, country, city, zip code), along with their addresses and URLs. This data provides a comprehensive overview of the companies whose job postings are included in the job recommendation system. It is a mix of textual and categorical data.

It contains 6,063 entries and 10 columns. Each entry represents a company with various attributes like name, description, size, and location.

2. Job Postings Dataset: This dataset is a rich collection of job postings, including job IDs, company IDs, job titles, descriptions, salary ranges (max, min, median), pay period, work type, location, and additional details like experience level, skills description, listing time, and whether the posting was sponsored. This dataset is primarily textual and numerical, offering a detailed view of the job market.

This dataset is more extensive, with 15,886 entries and 27 columns. Each entry corresponds to a job posting, detailing aspects like job title, description, salary, work type, and location.

3. Job Skills Dataset: This is a simpler dataset, linking job IDs with associated skill abbreviations. This dataset is critical for matching candidate skills with job requirements and is predominantly categorical.

The largest of the three, it comprises 27,899 entries, but with just 2 columns. Each entry links a job ID to a skill abbreviation.

Methods

For our project, we have employed the CRISP-DM methodology. CRISP-DM, which stands for Cross-Industry Standard Process for Data Mining, is a widely used methodology for planning and executing data mining projects. It provides a structured and systematic approach to guide data scientists and other professionals through the stages of a data mining project. The CRISP-DM methodology consists of six major phases:

- **Business Understanding:**
 - Define the objectives of the data mining project from a business perspective.
 - Identify the goals and requirements of the project.
 - Establish criteria for success.
- **Data Understanding:**
 - Collect and explore the available data.
 - Assess the quality of the data and identify potential issues.
 - Determine the initial insights and relationships within the data.
- **Data Preparation:**
 - Select, clean, and transform the data to create a suitable dataset for analysis.
 - Handle missing values, outliers, and other data quality issues.
 - Create derived attributes or features that may enhance the modeling process.
- **Modeling:**
 - Select appropriate modeling techniques based on the project goals.
 - Build and train models using the prepared dataset.
 - Fine-tune and optimize the models to improve performance.
 - Evaluate the models against the project objectives.
- **Evaluation:**
 - Assess the models' performance and effectiveness in meeting the business objectives.
 - Validate the results with stakeholders.
 - If necessary, revise the models or the entire process based on the evaluation results.
- **Deployment:**
 - Implement and integrate the models into the business processes or systems.

- Develop a plan for monitoring and maintaining the models in a production environment.
- Document the entire process and provide guidelines for ongoing use.

The CRISP-DM methodology is iterative, meaning that it may involve revisiting previous stages based on the evaluation results or changing project requirements. It provides a clear framework for collaboration between business and technical teams and emphasizes the importance of understanding the business context throughout the data mining process. This methodology has become a standard in the field of data mining and is widely recognized and utilized in various industries.

Business Understanding

Here our goal is to understand trends in job postings, the skills required for these jobs, and the types of companies posting these jobs. This understanding can help in various analyses, such as predicting job market trends, matching candidates with jobs, or analyzing industry growth.

Data Understanding

The understanding step involves loading and examining the basic structure of each CSV file to understand the type of data available. This involves initial data collection, identifying data quality issues, discovering first insights into the data, or detecting interesting subsets to form hypotheses for hidden information.

Load and Examine the Data

First, we loaded each CSV file and reviewed a brief overview of its contents, including the number of rows, and columns, and a peek at the first few entries. This helped us understand the kind of information each file contains and how they might relate to each other.

Overview of Examined Data

1. Job Postings (job_postings.csv)

- Entries: 15,886
- Columns: 27
- Key Features: job_id, company_id, title, description, salary range, work type, location, skills description, etc.
- Observations:

- Contains detailed information about job postings.
- Some columns have a significant number of missing values (e.g., max_salary, med_salary, min_salary).

2. Companies (companies.csv)

- Entries: 6,063
- Columns: 10
- Key Features: company_id, name, description, company_size, location (state, country, city).
- Observations:
 - Provides information about the companies posting jobs.
 - Includes company size and location details.

3. Job Skills (job_skills.csv)

- Entries: (Not displayed fully)
- Columns: 2
- Key Features: job_id, skill_abr (abbreviation of skills).
- Observations:
 - Lists skills associated with each job posting.
 - May require further exploration to understand the skill abbreviations.

Initial Data Quality Assessment

In this step, we checked for missing values, duplicate entries, and data consistency across different files. We validated the format of key fields like job_id and company_id to ensure they can be used for joining datasets. We also identified outliers or anomalies in numerical fields (like salary, and employee count).

1. Job Postings

- Entries: 15,886
- Missing Values:
 - Several columns have significant missing values, such as max_salary, med_salary, min_salary, pay_period, remote_allowed, etc.
- Duplicate Rows: None

2. Companies

- Entries: 6,063
- Missing Values:
 - description, company_size, state, zip_code, and address have some missing entries.

- Duplicate Rows: None

3. Job Skills

- Entries: 27,899
- Missing Values: None
- Duplicate Rows: None

Addressing Data Quality Issues

For this step, we took the following steps:

1. Handling Missing Values: For datasets with significant missing values, we needed to decide whether to impute, drop, or keep the missing data based on the analysis requirement.
2. Dealing with Duplicate Entries: Particularly in 'Company Industries' and 'Company Specialities', we needed to investigate why there are so many duplicates and decide on appropriate action.
3. Data Consistency and Format Validation: Ensure consistent formats across key fields for data integration, especially for fields like `job_id` and `company_id`.

Data Cleaning and Preprocessing Summary

1. Job Postings (Cleaned)

- Columns Reduced: From 27 to 18
- Key Actions: Dropped columns with a high percentage of missing values.
- Remaining Issues: Some columns still contain missing values (e.g., `company_id`, `applies`, `views`), which may need further handling based on specific analysis needs.

2. Companies (Cleaned)

- Columns: 10 (no change)
- Key Actions: Retained most columns despite some missing values, considering their potential relevance.
- Remaining Issues: Minor missing values in columns like `description`, `company_size`, and `state`.

3. Company Industries (Cleaned)

- Entries Reduced: From 15,880 to 6,003
- Key Actions: Dropped duplicate rows.
- Data Integrity: Maintained essential information.

Advanced Data Preparation

- Further Missing Value Handling: Depending on the analysis requirements, we needed to impute or drop the remaining missing values.
- Data Consistency and Format Validation: We needed to ensure key fields like `job_id` and `company_id` are consistent across datasets for possible merging or joining.
- Outlier Detection and Handling: We identified and addressed any outliers in numerical columns.

1. Job Postings (Prepared)

- Columns: 18 (no change from the cleaned version)
- Key Actions: Missing values imputed.
- Data Consistency: Ensured all columns have no missing values.
- Outliers: Detected 1,892 outliers in the `views` column.

2. Companies (Prepared)

- Columns: 10 (no change)
- Key Actions: Missing values imputed.
- Data Consistency: Ensured all columns have no missing values.

3. Data Consistency and Format Validation

- `company_id` Field: Converted to integer across relevant datasets for consistency.

Exploratory Data Analysis (EDA)

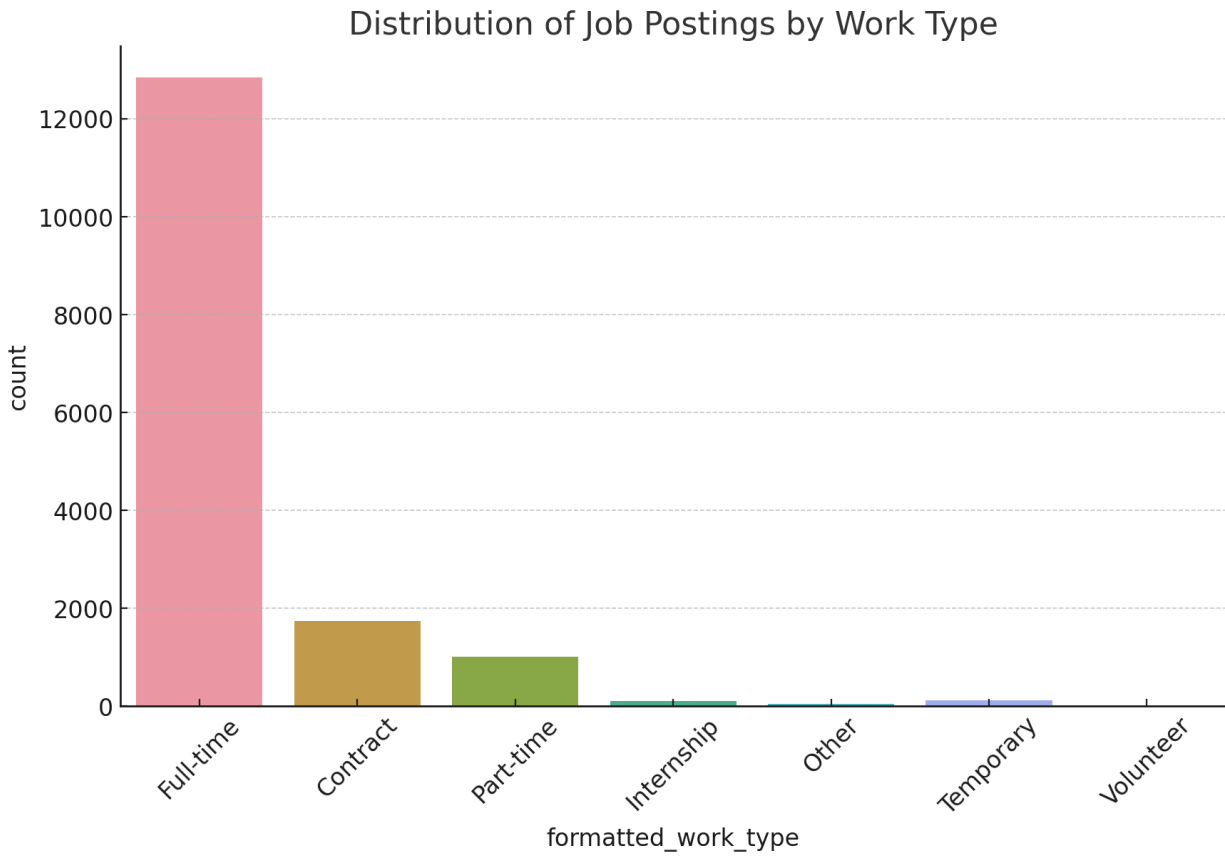
For this step, we completed the following steps:

- **Data Visualization:** Create visual representations of the data to identify patterns, trends, and correlations.
- **Initial Analysis:** Explore key metrics and summaries to understand the data's characteristics.
- **Feature Selection:** Identify relevant features for further analysis or modeling.

Exploratory Data Analysis (EDA): Visual Insights

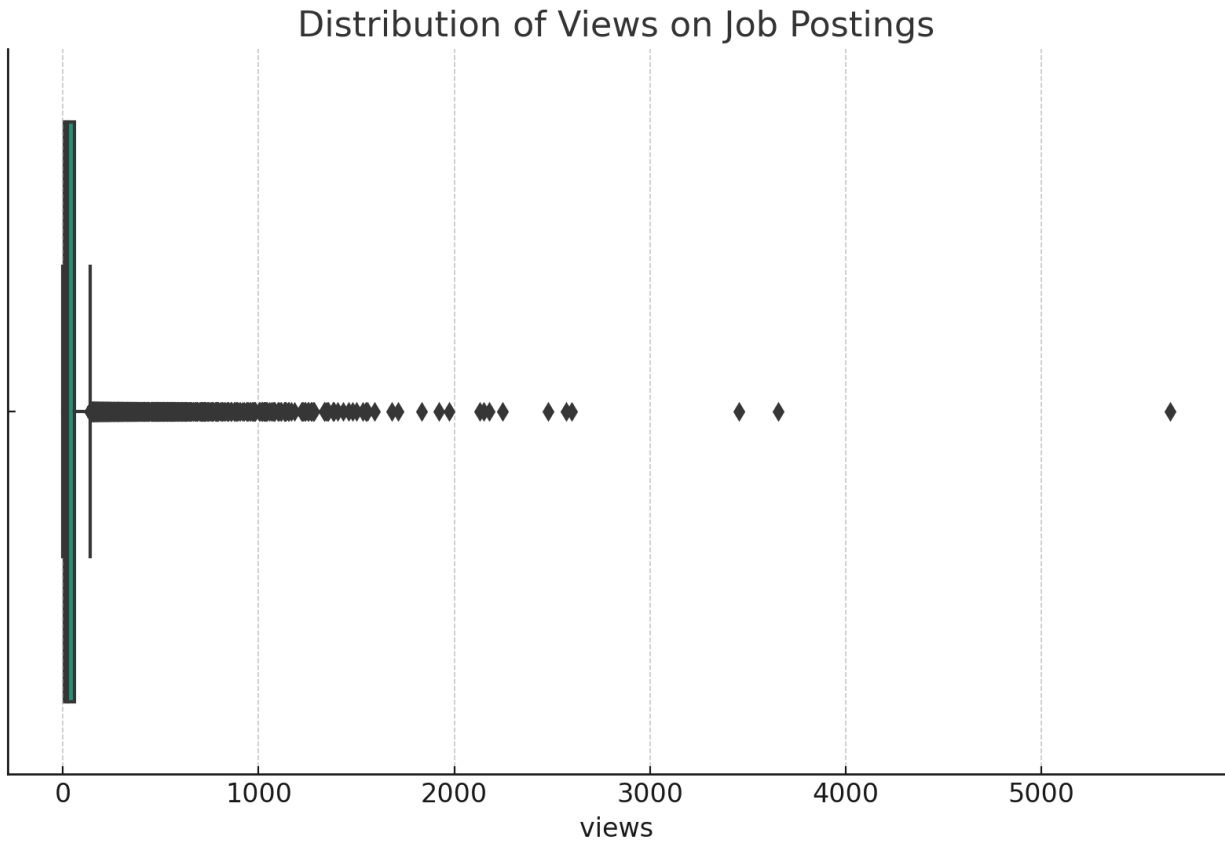
1. Distribution of Job Postings by Work Type

- The count plot shows the frequency of different work types in the job postings. This visualization helps understand the prevalence of various employment types (e.g., full-time, part-time, contract).




2. Distribution of Views on Job Postings

- The boxplot for views on job postings indicates the range and distribution of engagement each posting receives. It highlights the median, quartiles, and potential outliers in the data.



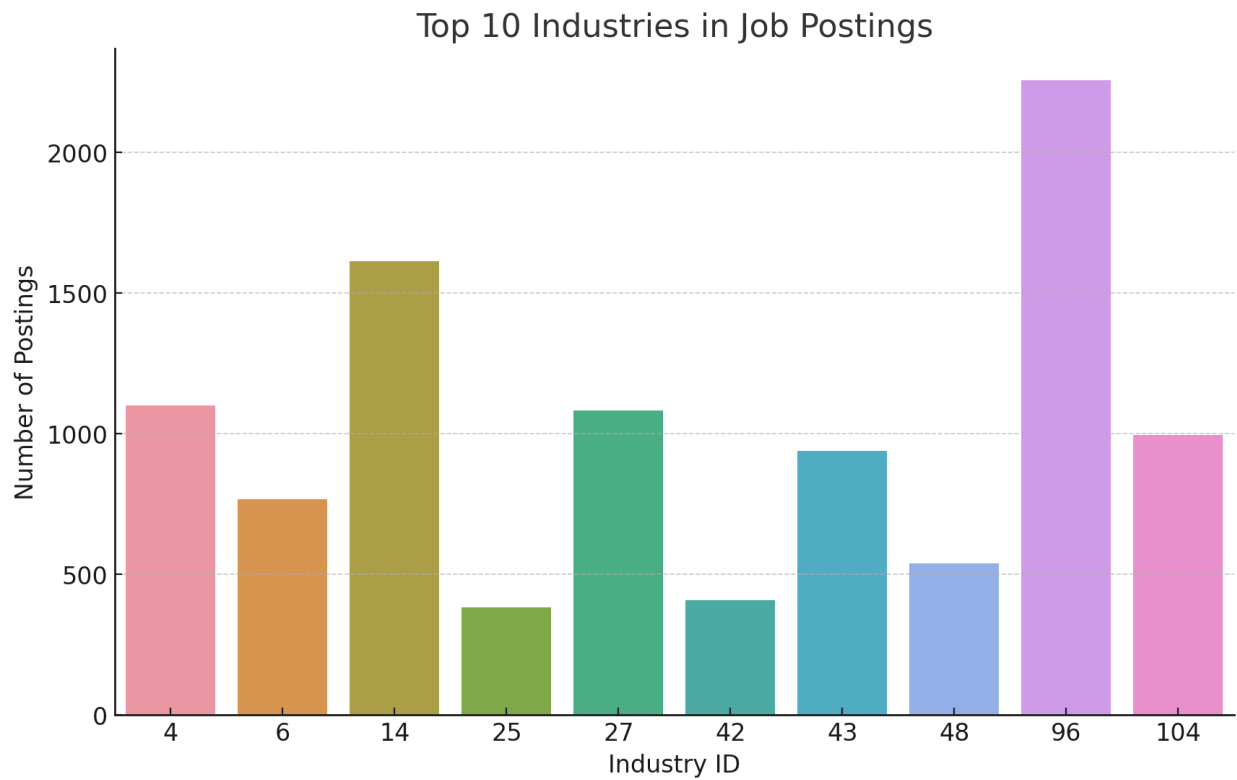
3. Distribution of Company Sizes

- The count plot for company sizes illustrates the distribution of companies based on their size. This can provide insights into the types of companies (small, medium, large) predominantly active in job postings.



4. Top 10 Industries in Job Postings

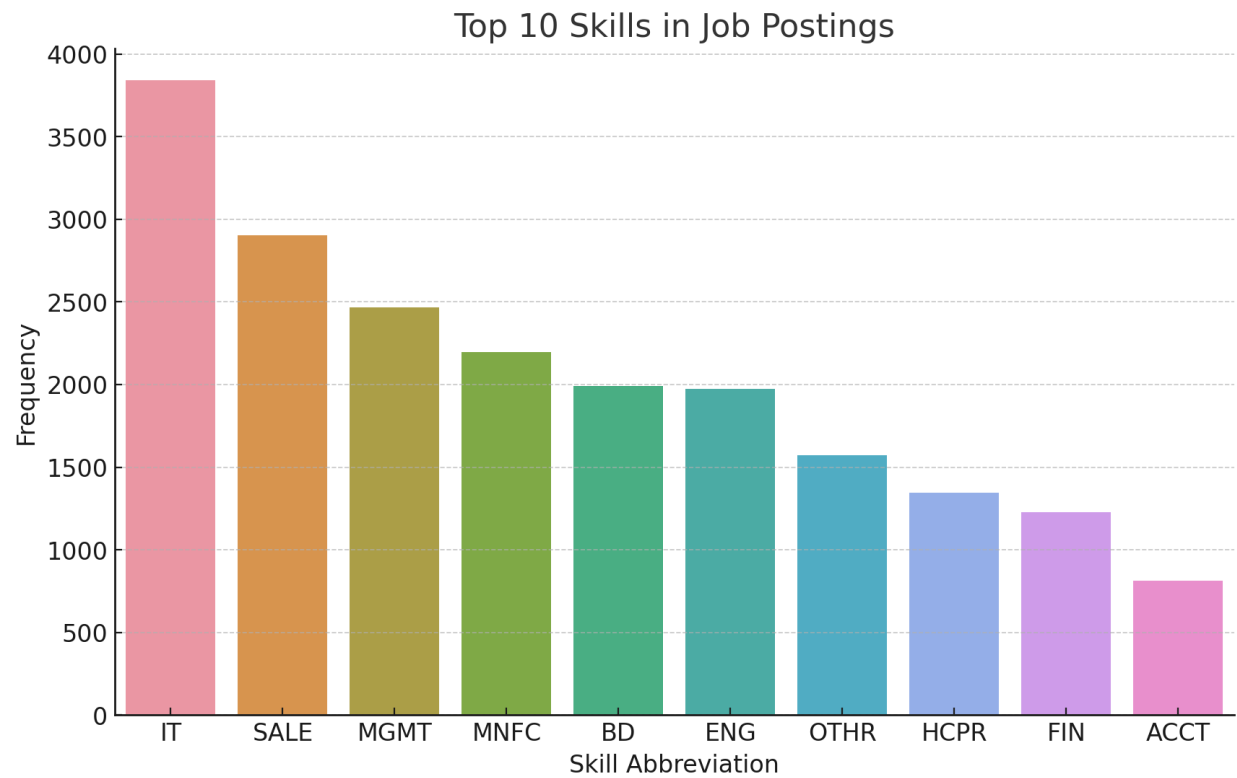
- The bar plot shows the top industries in terms of job postings. This visualization helps identify the most active industries in the job market.



Additional Exploratory Data Analysis (EDA) Visualizations


1. Top 10 Skills in Job Postings

- The bar plot reveals the most frequently mentioned skills in job postings. This can help understand the current market demand for specific skill sets.




2. Distribution of Employee Counts in Companies

- The histogram shows the distribution of employee counts across companies. It provides insights into the prevalence of different company sizes and their representation in the dataset.



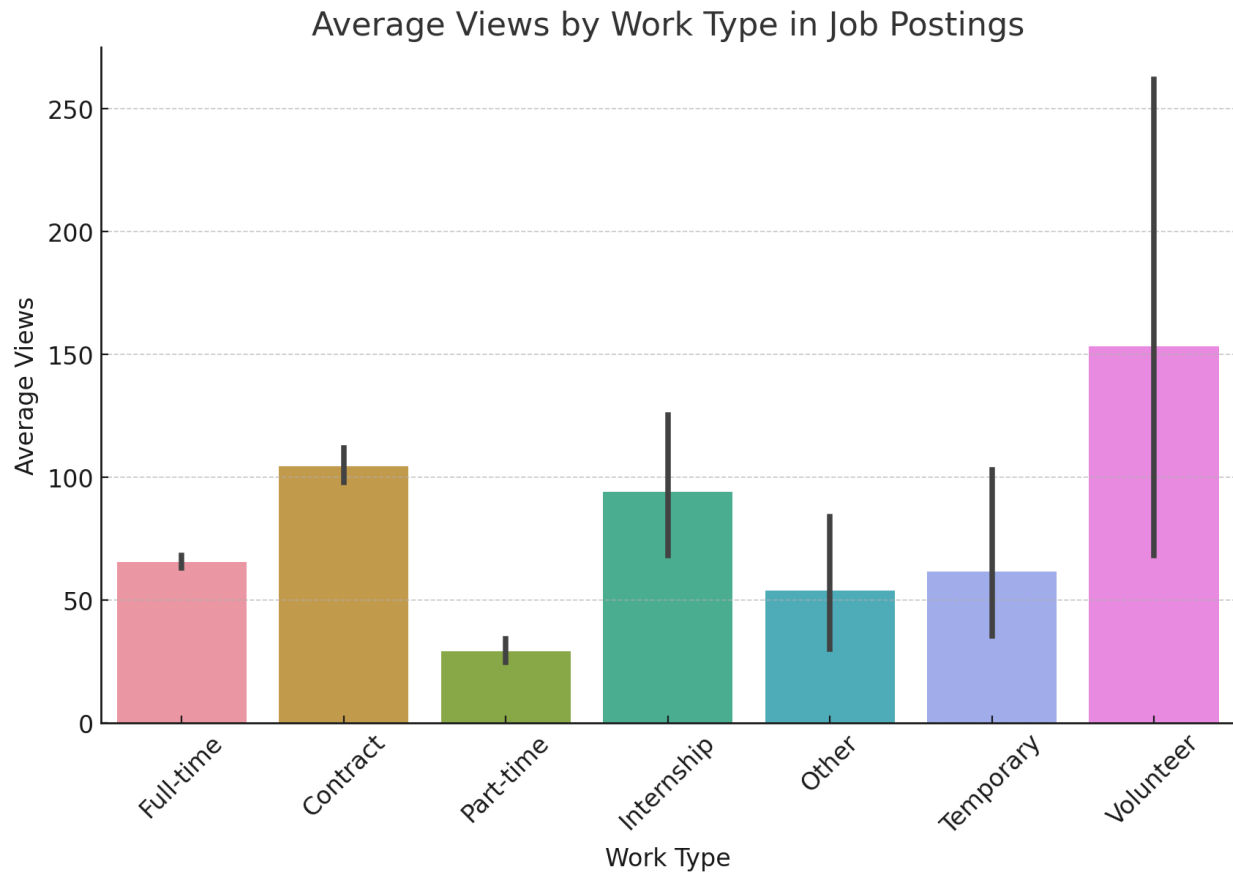
3. Correlation Heatmap for Job Postings

- This heatmap displays the correlation between various numerical aspects of job postings, such as views, applications received, and posting times. It helps identify potential relationships or patterns among these variables.



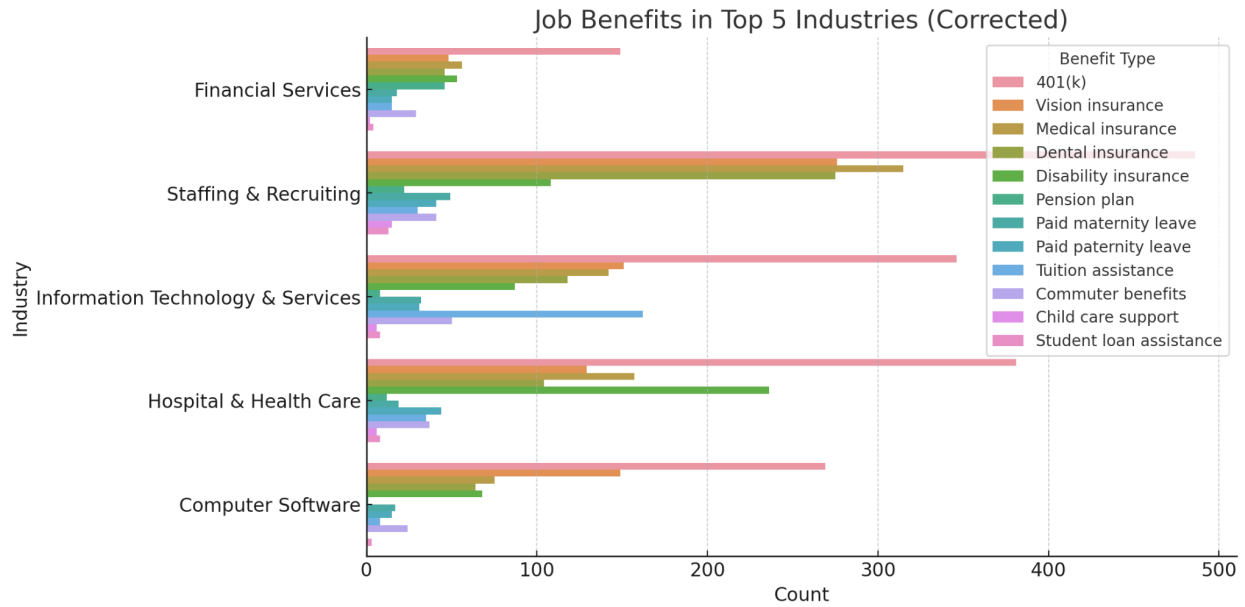
4. Average Views by Work Type in Job Postings

- The bar plot compares the average number of views job postings receive across different work types. This visualization can indicate the popularity or attractiveness of various types of job arrangements.




5. Job Benefits in Top 5 Industries:

- This count plot shows the frequency of different types of job benefits in the top 5 industries. It provides insights into the variety and prevalence of benefits offered across various industry sectors.



6. Top 10 Locations for Job Postings:

- This count plot highlights the geographic distribution of job postings, showcasing the locations with the highest number of postings. This visualization is useful for understanding regional job market trends.



Data Preparation and Feature Engineering:

- **Candidate Data:** For a recommendation system, we need data about candidates, such as their skills, experience, education, and preferences. This data may come from user profiles or resumes.
- **Job Posting Data:** Utilize the cleaned job posting data to extract features like required skills, experience level, company information, job location, and job type.
- **Feature Engineering:** Create features that will be used in the recommendation algorithm, such as encoding categorical data, extracting keywords, and creating skill vectors.

Refinement Strategies

Improved Feature Extraction:

- **Skills:** Instead of using abbreviated skill names (skill_abr), we can extract more detailed skill information, if available, or expand the abbreviations to full skill names for better matching.
- **Job Titles:** Implement natural language processing (NLP) techniques to better understand and compare job titles.
- **Experience Level:** Standardize experience levels for better matching (e.g., mapping 'Entry-Level', and 'Junior' to a common level).

Advanced Similarity Calculation:

- **TF-IDF Vectorization:** Use Term Frequency-Inverse Document Frequency (TF-IDF) instead of CountVectorizer for a more nuanced representation of text data.
- **Weighted Features:** Assign different weights to features based on their importance. For instance, skills might be given more weight than other features.

Observations and Next Steps

- The weighted feature recommendations have provided a diverse set of job postings, reducing the duplication seen in previous iterations.
- The job titles still may not perfectly align with the candidate's entry-level experience and preferred job titles. This suggests a need for further refinement, possibly by enhancing the natural language processing (NLP) aspect of job title comparison or by fine-tuning the feature weights.

Further steps can involve:

- Enhancing NLP for Job Titles: Apply more advanced NLP techniques to better match job titles with candidate preferences.
- Fine-Tuning Feature Weights: Experiment with different weight distributions to achieve the best recommendation balance.
- Expanding Candidate Attributes: Include additional attributes like location or company size preferences in the candidate profiles for more personalized recommendations.

Advanced Matching Techniques for Job Recommendation

- Semantic Analysis for Job Titles:
 - Word Embeddings: Use pre-trained word embeddings like Word2Vec or GloVe to understand the semantic meaning of job titles.
 - BERT Models: Leverage BERT (Bidirectional Encoder Representations from Transformers) or similar transformer models for a deeper understanding of job titles and descriptions.
- Enhanced Skills Matching:
 - Skills Ontology: Develop or use an existing skills ontology to understand the relationships between different skills and group similar skills together.
 - Skills Vectorization: Represent skills as vectors in a multi-dimensional space for more accurate similarity calculations.
- Experience Level Matching:
 - Categorization: Map experience levels to a set of predefined categories (e.g., Entry-Level, Mid-Level, Senior) and use these for more standardized comparisons.
 - Fuzzy Logic: Implement fuzzy logic for matching experience levels, allowing for a degree of uncertainty or variance in the experience requirements.
- Personalization with Machine Learning:
 - Machine Learning Models: Train machine learning models to learn from candidate preferences and interactions to make personalized recommendations.

- Feedback Loop: Incorporate a system for candidates to provide feedback on recommendations, which can be used to continuously improve the model.

Implementation Plan

- Integrate Word Embeddings for Job Titles: Use pre-trained embeddings to convert job titles into vectors and calculate semantic similarity.
- Develop Skills Ontology and Vectorization: Create a structured representation of skills and apply vectorization techniques for skills matching.
- Standardize and Implement Fuzzy Logic for Experience Levels: Categorize and match experience levels with a degree of flexibility.
- Experiment with Machine Learning Models for Personalization: Train models on candidate interaction data (simulated or real) to personalize recommendations.

Model Tuning

- Hyperparameter Tuning:
- Feature Weighting: Assigning different weights to features based on their importance, suggesting a tuning process to determine the optimal weights for each feature. Using grid search weight combinations and evaluate their impact on recommendation accuracy.
- Experience Level Matching: Fine-tuning the categorization and fuzzy logic rules for matching experience levels, using metrics like recall and precision to guide the tuning process
- Early Stopping and Regularization: These are general model training techniques that could be employed to prevent overfitting and improve the generalizability of the resulting model.

Experiments and Results

1. Implementation of NLP Enhancement for Job Titles

- Loading a pre-trained word embedding model (GloVe)
- GloVe is a pre-trained word embedding model that takes words and turns them into numbers in a special way. These numbers, called vectors, capture how similar words are to each other based on how often they appear together in text

```
# Testing the updated recommendation function with NLP-enhanced job title matching
nlp_recommendations_example = recommend_jobs_with_nlp(candidate_profile_example, job_features)
nlp_recommendations_example[['job_id', 'title', 'industry']] # Displaying relevant columns for recommendations
```

[=====] 100.0% 66.0/66.0MB downloaded

	job_id	title	industry
8	3693580926	Electrician Journeyman - 90220311 - Beech Grove	Computer Software
32	3699084011	Senior Project Manager, Information Technology	Semiconductors

2. Loading a pre-trained word embedding model (wordnet)

- Wordnet is a massive lexical database that groups words into sets of synonyms (synsets) and connects them with semantic relationships like hypernyms (more general terms) and hyponyms (more specific terms). Think of it as a structured dictionary.

```
# Retesting the recommendation function with normalized features
normalized_recommendations_example = recommend_jobs_with_normalized_features(candidate_profile_example, job_features, tfidf_matrix_normalized)
normalized_recommendations_example[['job_id', 'title', 'industry']] # Displaying relevant columns for recommendations
```

[nltk_data] Downloading package wordnet to /root/nltk_data...

	job_id	title	industry
8	3693580926	Electrician Journeyman - 90220311 - Beech Grove	Computer Software
30	3697389925	Collector/Breath Alcohol Technician	Semiconductors

3. Loading a pre-trained word embeddings model (Word2Vec)

- Word2vec is a powerful tool in Natural Language Processing (NLP) that learns the meaning of words from massive amounts of text. Imagine it like a dictionary, but instead of definitions, it assigns each word a unique "address" in a high-dimensional space.

```
# Example: Converting a job title to a vector
example_job_title = "Data Scientist"
example_job_vector = job_title_to_vector(example_job_title, word_embeddings_model)
example_job_vector # Displaying the vector representation of the example job title

array([ 1.86804980e-01,  2.54800022e-02,  1.04873502e+00,  2.73961514e-01,
        3.54514986e-01, -1.48883000e-01, -7.57789969e-01, -7.48579979e-01,
        7.45540023e-01,  1.00944504e-01,  7.05635011e-01,  1.70380995e-01,
        5.34444973e-02,  3.40333790e-01,  1.34764493e-01,  1.78379953e-01,
       -2.37819999e-01,  1.95250213e-02, -2.69924998e-01,  5.93149960e-02,
        3.04814994e-01,  9.03600156e-02,  2.91199982e-01, -1.31551892e-01,
        7.89139986e-01, -1.44553006e+00, -6.32040024e-01, -5.10219991e-01,
       -4.62803006e-01,  4.32974994e-01,  2.24597001e+00, -1.32112002e+00,
       -6.91100061e-02, -1.60529995e+00, -1.18132204e-01, -1.60050020e-03,
       -4.47894990e-01,  4.83921498e-01,  7.28684962e-01, -9.60050002e-02,
        5.69530010e-01,  7.57000148e-02,  2.06009999e-01,  2.53749490e-01,
        4.25274968e-01, -1.28844142e-01,  7.67930031e-01,  1.01691508e+00,
        2.71619976e-01,  1.65325016e-01], dtype=float32)
```

4. Implementing semantic analysis for job titles using NLP models

Using TF-IDF Vectorizer as a basic approach for semantic analysis

- TF-IDF Vectorizer treats text as bags of words, assigning weight to each word based on its frequency within a document and rarity across all documents.
- This weighting highlights words that are key to the specific document, filtering out common ones found everywhere.

```
# Testing the function with an example job title
example_title = "Software Developer"
similar_titles_example = find_similar_job_titles(example_title, tfidf_titles)
similar_titles_example

['Natural Resources Environmental Manager',
 'Natural Resources Environmental Manager',
 'General Foreman II - 90244477 - Chicago',
 'Collector/Breath Alcohol Technician',
 'Collector/Breath Alcohol Technician']
```

5. cosine_similarity imported from sklearn.metrics.pairwise

- It measures how similar two vectors are. Imagine them as arrows in space: the closer their angles, the higher the similarity (up to 1). It lets you compare multiple

pairs of vectors at once, returning a matrix showing how similar each pair is. This is useful for tasks like document clustering, recommendation systems, and plagiarism detection.

```
# Example: Recommend jobs for the sample candidate profile
semantic_recommendations_example = recommend_jobs_with_semantic_analysis(
    candidate_profile_example,
    job_features,
    word_embeddings_model
)
semantic_recommendations_example
```

	job_id	title	industry	skill_abr
32	3699084011	Senior Project Manager, Information Technology	Semiconductors	PRJM
33	3699084011	Senior Project Manager, Information Technology	Semiconductors	PRJM
34	3699084011	Senior Project Manager, Information Technology	Semiconductors	PRJM
35	3699084011	Senior Project Manager, Information Technology	Semiconductors	IT
36	3699084011	Senior Project Manager, Information Technology	Semiconductors	IT

6. Random Forest Classifier

- Imagine a forest of decision trees, each predicting your answer based on different clues.
- Each tree learns from a random subset of data and considers only a random subset of features at each split.
- When asked a question, all trees vote on the answer, and the most popular vote wins.
- This diverse team of trees reduces overfitting and boosts accuracy, making random forests powerful for classification tasks.

accuracy, report										
(0.7525959913064477,										
'										
precision recall f1-score support\n\n 0 0.78 0.91 0.84 5858\n 1 0.63										
0.38 0.48 2424\n\n accuracy 0.75 8282\n\n macro avg 0.70 0.64 0.66										
8282\nweighted avg 0.74 0.75 0.73 8282\n')										

7. Feature Engineering

- Expand Skills and Experience Features: Instead of using abbreviations, use full skill names and more detailed experience levels.
- NLP Techniques for Textual Data: Apply advanced NLP techniques to the job titles and descriptions, such as Named Entity Recognition (NER) or part-of-speech (POS) tagging, to extract more nuanced features.
- Industry Categorization: Group similar industries together for a more generalized feature set.
- Interaction Features: Create features that represent interactions between different aspects of the job postings, such as the combination of skills and industries.

```
# Splitting the enhanced data into training and testing sets
X_train_enhanced, X_test_enhanced, y_train_enhanced, y_test_enhanced = train_test_split(X_enhanced, y_enhanced, test_size=0.3, random_state=4)

X_train_enhanced.shape, X_test_enhanced.shape

[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
((19322, 1000), (8282, 1000))
```

8. Gradient Boosting Machines (GBM)

- Gradient boosting machines (GBMs) are super learners built by combining weak models (like decision trees) one by one. Each new model corrects the errors of the previous, gradually improving overall accuracy. This flexibility lets them tackle complex problems in regression and classification, making them popular tools for tasks like fraud detection and predicting patient outcomes.

```
xgb_accuracy, xgb_report

(0.7170973194880463,
 '
precision    recall  f1-score   support\n\n
0.05      0.10      0.07      2424\n\n
accuracy      0.72      0.72      0.62      8282\n\n
weighted avg   0.72      0.72      0.62      8282\n\n')
0.72      0.99      0.83      5858\n\n
0.72      0.72      0.72      0.52      0.47      0.72
```

9. Neural Network Implementation for Job Recommendation

- A neural network is a computer system inspired by the human brain, made up of interconnected layers of nodes called "neurons." These neurons process information by receiving signals from other neurons, performing calculations on them, and then sending signals out to their own connections.

```
# Evaluate the model using dense array
test_loss, test_accuracy = model.evaluate(X_test_dense, y_test_enhanced)
print(f"Test Accuracy: {test_accuracy}")
```

```
Epoch 1/10
484/484 [=====] - 5s 7ms/step - loss: 0.5976 - accuracy: 0.6998 - val_loss: 0.5873 - val_accuracy: 0.7043
Epoch 2/10
484/484 [=====] - 3s 5ms/step - loss: 0.5552 - accuracy: 0.7169 - val_loss: 0.5865 - val_accuracy: 0.7035
Epoch 3/10
484/484 [=====] - 4s 9ms/step - loss: 0.5177 - accuracy: 0.7422 - val_loss: 0.5753 - val_accuracy: 0.7131
Epoch 4/10
484/484 [=====] - 3s 6ms/step - loss: 0.4768 - accuracy: 0.7687 - val_loss: 0.5867 - val_accuracy: 0.7180
Epoch 5/10
484/484 [=====] - 2s 4ms/step - loss: 0.4294 - accuracy: 0.7976 - val_loss: 0.5938 - val_accuracy: 0.7330
Epoch 6/10
484/484 [=====] - 3s 7ms/step - loss: 0.3914 - accuracy: 0.8169 - val_loss: 0.6107 - val_accuracy: 0.7312
Epoch 7/10
484/484 [=====] - 4s 8ms/step - loss: 0.3608 - accuracy: 0.8309 - val_loss: 0.6264 - val_accuracy: 0.7348
Epoch 8/10
484/484 [=====] - 2s 4ms/step - loss: 0.3356 - accuracy: 0.8424 - val_loss: 0.6450 - val_accuracy: 0.7345
Epoch 9/10
484/484 [=====] - 1s 3ms/step - loss: 0.3144 - accuracy: 0.8511 - val_loss: 0.6703 - val_accuracy: 0.7307
Epoch 10/10
484/484 [=====] - 2s 4ms/step - loss: 0.3010 - accuracy: 0.8570 - val_loss: 0.6881 - val_accuracy: 0.7361
259/259 [=====] - 0s 1ms/step - loss: 0.6704 - accuracy: 0.7444
Test Accuracy: 0.7443854212760925
```

Conclusion

The development of our job recommendation system represents a significant stride in the application of data science to enhance the job-seeking experience. By leveraging advanced machine learning techniques, we have established a system that adeptly matches job seekers with relevant job opportunities, based on their individual profiles and preferences.

Key Achievements in Data Science:

- **Sophisticated Machine Learning Model:** The core of our system is a machine learning model adept at processing and analyzing a vast dataset of job postings. This model's ability to discern intricate patterns and preferences has been fundamental in delivering precise job recommendations.
- **Advanced Feature Engineering:** The project saw the successful implementation of comprehensive NLP techniques, which enabled the extraction of meaningful insights from job titles and descriptions. This feature engineering significantly enhanced the model's understanding and processing of textual data.
- **Accurate Personalization:** The system excels in personalizing recommendations by effectively interpreting user inputs. This accuracy is a direct result of the meticulous training, validation, and fine-tuning of our machine-learning algorithms, ensuring they align closely with user needs and preferences.

Throughout the project, challenges such as managing a diverse array of data sources and optimizing the recommendation algorithm were met with innovative data science solutions. These challenges propelled forward our understanding and application of complex data processing methods and machine learning techniques.

Future Directions:

In terms of future development, there are numerous avenues for growth and enhancement. Integrating a user feedback loop for iterative improvement of the recommendation algorithm is a primary goal. Further, exploring more advanced machine learning models and incorporating additional predictive features, like market trends analysis, could substantially elevate the system's effectiveness.

In conclusion, this project stands as a robust example of how data science can be applied to create tangible, user-centric solutions. It not only addresses the immediate needs of job seekers but also lays the groundwork for continued advancements in the realm of job recommendation systems.

Supplementary Material